

Statistical Calculation and Software

Assignment 4

Hanbin Liu 11912410

4.1

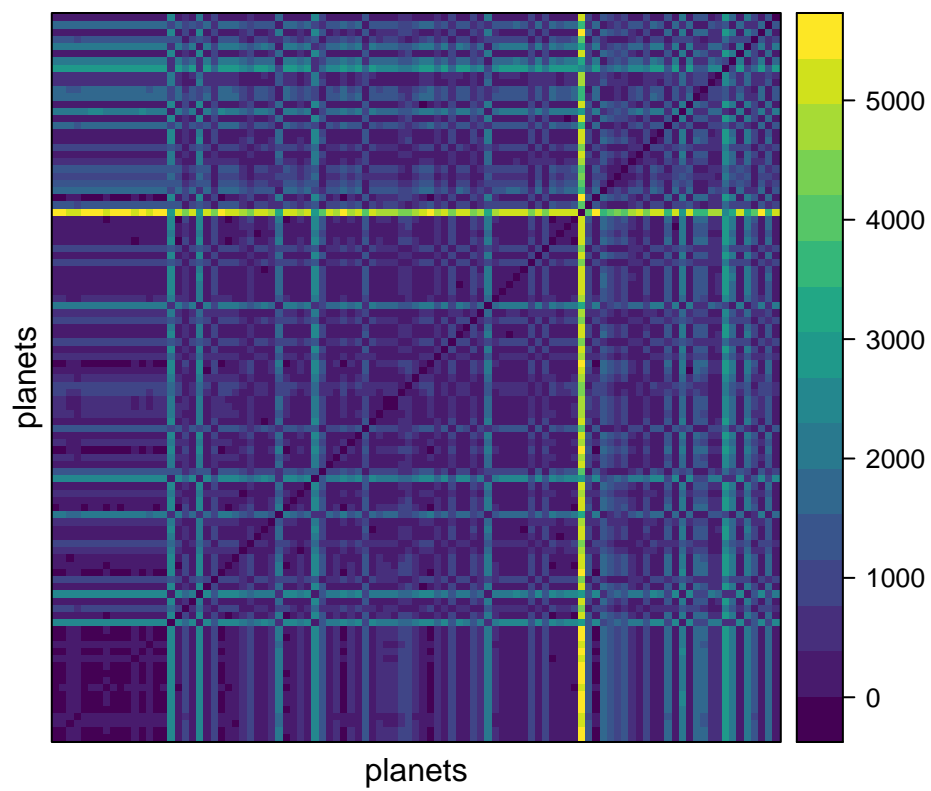
```
library(HSAUR3)
```

```
##      tools
```

```
data(planets)
```

(a)


```
library(lattice)
library(viridisLite)
planets_dist <- dist(planets)
levelplot(
  as.matrix(planets_dist),
  xlab = "planets",
  ylab = "planets",
  col.regions = viridis(100),
  scales = list(draw = FALSE)
)
```



```
planets_single <- hclust(planets_dist, method = "single")
planets_complete <- hclust(planets_dist, method = "complete")
planets_average <- hclust(planets_dist, method = "average")


plot(planets_single,
     main = "Single Linkage",
     sub = "",
     xlab = "",
)
```

Single Linkage




```
plot(planets_complete,  
     main = "Complete Linkage",  
     sub = "",  
     xlab = "")
```

Complete Linkage



```
plot(planets_average,  
     main = "Average Linkage",  
     sub = "",  
     xlab = "")
```

Average Linkage



```
planets_cluster_single <- cutree(planets_single, h = 1000)
planets_cluster_complete <- cutree(planets_complete, h = 1500)
planets_cluster_average <- cutree(planets_average, h = 900)
planets_cluster_single
```

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
## 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## 101
## 1
```

```
planets_cluster_complete
```

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 3 1
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
## 2 1 3 1 1 1 1 3 1 1 1 3 1 1 1 1 2 3 1 1
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
```

```
## 1 1 1 3 1 1 1 1 3 3 1 1 1 3 1 3 1 1 3 1
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
## 2 1 1 1 1 1 3 1 3 1 1 1 1 4 3 1 3 3 3 3
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
## 1 1 3 1 1 3 1 2 1 3 3 1 1 2 2 1 2 3 1 3
## 101
## 1
```

```
planets_cluster_average
```

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 3 1
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
## 2 1 3 1 1 1 1 3 1 1 1 3 1 1 1 1 2 3 1 1
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
## 1 1 1 3 1 1 1 1 3 3 1 1 1 3 1 3 1 1 3 1
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
## 2 1 1 1 1 1 3 1 3 1 1 1 1 4 3 1 3 3 3 3
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
## 1 1 3 1 1 3 1 2 1 3 3 1 1 2 2 1 2 3 1 3
## 101
## 1
```

(b)

```
## min-max standardized
rge <- apply(planets, 2, max) - apply(planets, 2, min)
planet.dat <-
  sweep(planets, 2, rge, FUN = "/") ### function = divide

## K=3
planet_kmeans3 <- kmeans(planet.dat, centers = 3)
planet_kmeans3
```

```
## K-means clustering with 3 clusters of sizes 14, 34, 53
```

```
##
```

```
## Cluster means:
```

```
##      mass      period      eccen
```

```
## 1 0.60560786 0.31606632 0.3953614
```

```
## 2 0.16777347 0.11500361 0.5343613
```

```
## 3 0.09576256 0.07984122 0.1315524
```

```
##
```

```
## Clustering vector:
```

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

```
## 3 3 2 3 3 3 3 2 3 3 3 3 3 3 3 3 3 3 3 2
```

```
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

```
## 2 3 2 3 3 3 3 3 2 2 2 2 3 2 3 3 3 3 2 2
```

```
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
```

```
## 3 3 2 3 3 2 2 2 2 3 3 3 3 2 3 3 3 3 3 3
```

```
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
```

```
## 3 3 3 2 2 3 3 2 2 3 2 2 3 1 2 3 2 2 2 2
```

```
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
## 2 3 3 3 2 1 2 1 2 1 1 2 1 1 1 1 1 1 1 1
## 101
## 1
##
## Within cluster sum of squares by cluster:
## [1] 1.719130 1.787017 1.755694
## (between_SS / total_SS = 57.2 %)
##
## Available components:
##
## [1] "cluster" "centers" "totss" "withinss" "tot.withinss"
## [6] "betweenss" "size" "iter" "ifault"
```

```
table(planet_kmeans3$cluster)
```

```
##
## 1 2 3
## 14 34 53
```

```
ccent <- function(cl) {
  f <- function(i)
    colMeans(planets[cl == i, ])
  x <- sapply(sort(unique(cl)), f)
  colnames(x) <- sort(unique(cl))
  return(x)
}

ccent(planet_kmeans3$cluster)
```

```
##           1           2           3
## mass      10.56786    2.9276471    1.6710566
## period 1693.17201 616.0760882 427.7105892
## eccen      0.36650    0.4953529    0.1219491
```

```
## K=5
planet_kmeans5 <- kmeans(planet.dat, centers = 5)
planet_kmeans5
```

```
## K-means clustering with 5 clusters of sizes 8, 17, 14, 30, 32
##
## Cluster means:
##      mass      period      eccen
## 1 0.11840974 0.44869904 0.2067152
## 2 0.21051744 0.12598649 0.6574656
## 3 0.61960704 0.24615398 0.4138542
## 4 0.09991595 0.03290963 0.0531931
## 5 0.09563037 0.07505714 0.3267260
##
## Clustering vector:
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## 4 4 5 5 4 4 4 5 4 4 4 5 4 5 4 4 1 5 4 2
```

```
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
## 2 4 5 4 4 4 4 5 5 5 2 1 4 5 5 5 1 1 5 5
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
## 4 4 2 1 5 5 2 5 2 5 5 4 4 2 5 5 5 5 4 4
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
## 1 5 5 5 5 5 5 2 2 4 5 2 4 1 2 4 1 5 2 2
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
## 2 4 4 4 2 3 3 3 2 3 3 2 3 3 3 3 3 3 3 3
## 101
## 3
##
## Within cluster sum of squares by cluster:
## [1] 0.5852326 0.7801019 1.1038061 0.4920874 0.5153905
## (between_SS / total_SS = 71.7 %)
##
## Available components:
##
## [1] "cluster" "centers" "totss" "withinss" "tot.withinss"
## [6] "betweenss" "size" "iter" "ifault"
```

```
table(planet_kmeans5$cluster)
```

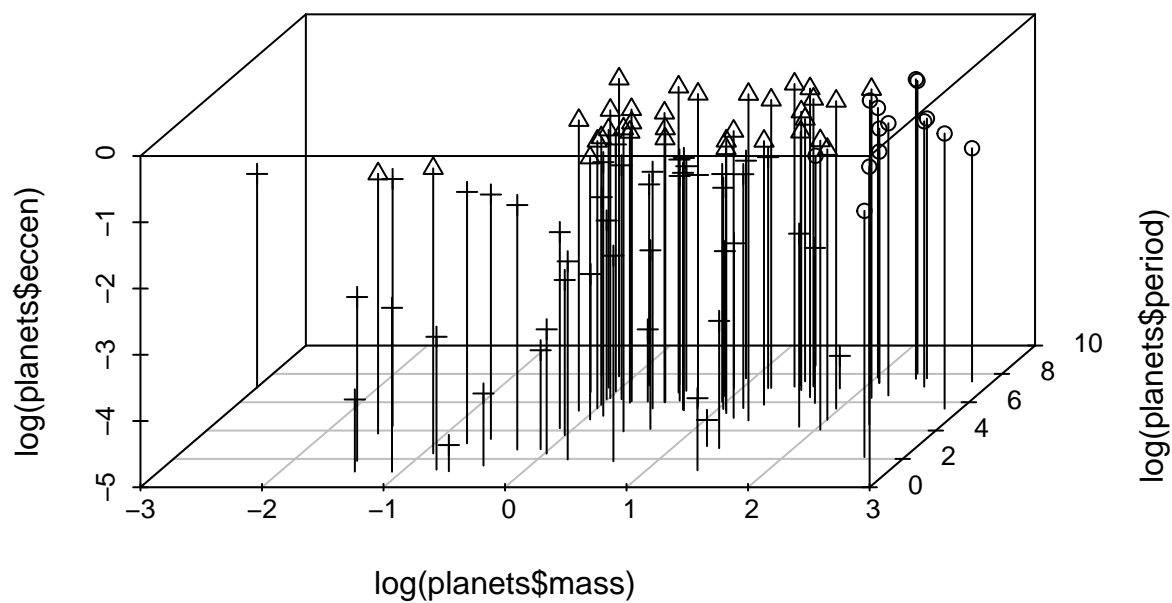
```
##
## 1 2 3 4 5
## 8 17 14 30 32
```

```
ccent(planet_kmeans5$cluster)
```

```
##           1           2           3           4           5
## mass      2.066250    3.6735294    10.8121429    1.743533    1.668750
## period 2403.687500  674.9115294 1318.6505856 176.297374 402.082219
## eccen     0.191625    0.6094706    0.3836429    0.049310    0.302875
```


- 3D scatterplot for K=3

```
library("scatterplot3d")
layout(matrix(1))
scatterplot3d(
  log(planets$mass),
  log(planets$period),
  log(planets$eccen),
  type = "h",
  angle = 55,
  scale.y = 0.7,
  pch = planet_kmeans3$cluster,
  y.ticklabs = seq(0, 10, by = 2),
  y.margin.add = 0.1,
)
```

- 3D scatterplot for K=5

```
library("scatterplot3d")
layout(matrix(1))
scatterplot3d(
  log(planets$mass),
  log(planets$period),
  log(planets$eccen),
  type = "h",
  angle = 55,
  scale.y = 0.7,
  pch = planet_kmeans5$cluster,
  y.ticklabs = seq(0, 10, by = 2),
  y.margin.add = 0.1,
)
```



(c)

```
logL <- function(param, x) {
  d1 <- dnorm(x, mean = param[2], sd = param[3])
  d2 <- dnorm(x, mean = param[4], sd = param[5])
  - sum(log(param[1] * d1 + (1 - param[1]) * d2))
}
x <- planets$eccen
startparam <-
  c(
    p = 0.5,
    mu1 = mean(x) / 2,
    sd1 = sd(x) / 2,
    mu2 = mean(x) * 2,
    sd2 = sd(x) * 2
  )

opp <-
  optim(
    startparam,
    logL,
    x = planets$eccen,
    method = "L-BFGS-B",
    lower = c(0.01, rep(0.01, 4)),
```

```

    upper = c(0.99, rep(1, 4))
  )
opp

```

```

## $par
##      p      mu1      sd1      mu2      sd2
## 0.18979358 0.02445790 0.02447048 0.34176842 0.18742605
##
## $value
## [1] -27.81359
##
## $counts
## function gradient
##      59      59
##
## $convergence
## [1] 0
##
## $message
## [1] "CONVERGENCE: REL_REDUCTION_OF_F <= FACTR*EPSMCH"

```

(d)

```

library("mclust")

```

```

## Package 'mclust' version 5.4.7
## Type 'citation("mclust")' for citing this R package in publications.

```

```

planet_mclust <- Mclust(planet.dat[3], G = 2)
print(planet_mclust)

```

```

## 'Mclust' model object: (V,2)
##
## Available components:
##   [1] "call"          "data"          "modelName"     "n"
##   [5] "d"             "G"             "BIC"           "loglik"
##   [9] "df"            "bic"           "ic1"           "hypvol"
##  [13] "parameters"    "z"             "classification" "uncertainty"

```

```

table(planet_mclust$classification)

```

```

##
##  1  2
## 21 80

```

```
## sample statistics
ind1 <- planet_mclust$classification == 1
ind2 <- planet_mclust$classification == 2
data.frame(mclust.BIC = c(
  p = length(x[ind1]) / length(x),
  mu1 = mean(x[ind1]),
  sd1 = sd(x[ind1]),
  mu2 = mean(x[ind2]),
  sd2 = sd(x[ind2])
))
```

```
##      mclust.BIC
## p    0.20792079
## mu1  0.02100000
## sd1  0.02110924
## mu2  0.34994125
## sd2  0.18283981
```

The results match the results in (c). Compare:

```
data.frame(loglikelihood = opp$par,
  mclust.BIC = c(21 / 101, mean(x[ind1]), sd(x[ind1]), mean(x[ind2]), sd(x[ind2])))
```

```
##      loglikelihood mclust.BIC
## p      0.18979358 0.20792079
## mu1     0.02445790 0.02100000
## sd1     0.02447048 0.02110924
## mu2     0.34176842 0.34994125
## sd2     0.18742605 0.18283981
```

(e)

```
planets_pca <- prcomp(planets, scale = TRUE)
```

coefficients for the first two principal components:

```
planets_pca$rotation[, 1]
```

```
##      mass    period    eccen
## 0.6423065 0.5229950 0.5602843
```

```
planets_pca$rotation[, 2]
```

```
##      mass    period    eccen
## -0.05996314 0.76306298 -0.64353657
```

score:

```

first.score <- predict(planets_pca)[, 1]
second.score <- predict(planets_pca)[, 2]
third.score <- predict(planets_pca)[, 3]

score <- data.frame(first = first.score,
                    second = second.score,
                    third = third.score)

score

```

```

##           first      second      third
## 1  -1.703521406  0.333412263 -0.317723645
## 2  -1.690670027  0.331302824 -0.334134464
## 3  -0.761204110 -0.670962589  0.521211858
## 4  -0.899957517 -0.460553244  0.384507035
## 5  -1.470968444  0.088828815 -0.142157505
## 6  -1.628866016  0.268597257 -0.296106234
## 7  -1.453818089  0.084051724 -0.166479222
## 8  -0.328364590 -1.185227634  0.857857639
## 9  -1.652030119  0.327266919 -0.380630115
## 10 -1.642874006  0.327080514 -0.390697571
## 11 -1.508768650  0.173731908 -0.269489321
## 12 -0.833611684 -0.572710156  0.356629745
## 13 -1.604774912  0.294443962 -0.381000325
## 14 -0.894526937 -0.475478120  0.269347580
## 15 -1.552872634  0.262983037 -0.384682573
## 16 -1.605775861  0.322962403 -0.435638183
## 17  0.223486270  2.278978687  0.921884194
## 18 -0.867342852 -0.492706556  0.218879861
## 19 -0.983648497  0.919756324 -0.003954419
## 20 -0.006926677 -1.137171697  0.953613065
## 21  1.526980544  0.685274524  2.100008258
## 22 -1.188140002 -0.023619320 -0.130265394
## 23 -0.013888070  0.301724319  0.827731460
## 24 -1.563279363  0.318593381 -0.486688478
## 25 -1.476672551  0.237173332 -0.407496711
## 26 -1.344904490  0.136129674 -0.310136647
## 27 -0.989331825  0.447137596 -0.035646395
## 28 -0.288827708  0.222501340  0.570830607
## 29 -0.341819216 -0.550128939  0.583362494
## 30 -0.385041510 -0.613191358  0.551978404
## 31  0.153460192 -0.799860208  1.015739803
## 32  0.673676747  0.790275183  1.328250468
## 33 -1.426139647  0.189555649 -0.393283377
## 34 -0.542969893 -0.649376687  0.404283622
## 35 -0.621790925  0.047784171  0.258425689
## 36 -0.614183042 -0.497909017  0.304088800
## 37  0.037908210  2.595166702  0.602398503
## 38 -0.347228837  1.041922025  0.372095526
## 39 -0.214558142 -0.567990814  0.606402616
## 40 -0.055057819 -0.681578440  0.749393536
## 41 -1.143520921 -0.092419140 -0.226744783
## 42 -1.171509632  0.618929291 -0.355165237
## 43  0.297759588 -1.629074895  1.010101238

```

```

## 44 -0.059345034 0.956988939 0.503417867
## 45 -0.658610158 -0.199814272 0.083340902
## 46 -0.087854829 -0.558042165 0.573455105
## 47 0.209516225 -0.958005671 0.851258877
## 48 -0.243206505 -0.443819721 0.426790238
## 49 0.848768694 -1.075064375 1.383582089
## 50 -0.348565080 0.480827967 0.226549467
## 51 -0.605104445 0.087566988 0.016871857
## 52 -1.095466131 0.048467670 -0.413595699
## 53 -0.993684420 -0.152041949 -0.315857932
## 54 0.711359076 -0.943966840 1.147558254
## 55 -0.587782978 -0.223802437 -0.022359296
## 56 -0.597372288 0.788846400 0.601262625
## 57 -0.420207682 -0.430700452 0.124074829
## 58 -0.565848109 -0.104750398 -0.080822730
## 59 -0.472965629 1.057663456 -0.200638924
## 60 -0.749298605 0.654486666 -0.401286743
## 61 0.503535941 1.648568524 0.570287524
## 62 -0.462877062 0.009093324 -0.141872980
## 63 -0.577216620 -0.076100630 -0.305972798
## 64 -0.147121897 -0.583840779 0.059855439
## 65 0.002211130 -0.805905002 0.170810635
## 66 -0.377211666 -0.191366341 -0.299477970
## 67 0.100494738 0.564043567 0.028580979
## 68 0.297972212 -1.166022173 0.326794165
## 69 0.919661797 -0.340259408 0.759698104
## 70 -1.070455711 0.258595779 -0.999317069
## 71 0.097165709 -0.532369568 -0.048184675
## 72 1.482287080 -2.462896010 1.235655294
## 73 -0.897754987 0.139323959 -1.004792716
## 74 2.603895761 4.458750841 1.599748890
## 75 1.442220277 -0.662437726 0.982456118
## 76 -0.949843973 0.205475045 -1.103550893
## 77 1.011203195 0.725098846 0.457916064
## 78 0.627835835 0.458158282 0.149491677
## 79 0.822348900 -0.284580296 0.296616717
## 80 1.218537242 -0.237997345 0.519327049
## 81 1.407137925 -1.416349503 0.635219323
## 82 -0.317535024 0.359110502 -1.138782140
## 83 0.005677719 1.136378075 -1.063202584
## 84 -0.256459179 -0.088240546 -1.262835911
## 85 1.536314165 -1.697724708 0.058291570
## 86 1.778750857 0.133980423 0.040114637
## 87 0.699735796 -0.908224948 -0.795792285
## 88 2.007260665 1.012444470 0.124150192
## 89 1.056167995 -1.357496410 -0.560477228
## 90 1.213952042 0.945137698 -0.697437474
## 91 1.435238879 0.603406441 -0.502383447
## 92 1.995773521 -1.494857341 -0.094099930
## 93 1.445323958 -0.507323990 -1.011218948
## 94 3.318745016 1.105864335 0.330322879
## 95 2.995113271 0.117878297 0.004245738
## 96 1.112024683 -0.780876060 -1.717950563
## 97 2.457878305 1.021389082 -0.848946457

```

```
## 98 2.069636084 0.062934791 -1.341881538
## 99 1.526318501 -0.741254008 -2.595981508
## 100 2.868605488 0.879282917 -2.483757420
## 101 2.619234399 -1.039237215 -2.756168649
```

(f)

```
## min-max standardized
rge <- apply(score, 2, max) - apply(score, 2, min)
score.dat <-
  sweep(score, 2, rge, FUN = "/") ### function = divide

## K=3
score_kmeans3 <- kmeans(score.dat[c(1, 2)], centers = 3)
```

compare:

```
score_kmeans3
```

```
## K-means clustering with 3 clusters of sizes 26, 30, 45
##
## Cluster means:
##      first      second
## 1 0.34247011 -0.013041092
## 2 -0.25256877 0.022724638
## 3 -0.02949244 -0.007614905
##
## Clustering vector:
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  2  2  3  2  2  2  2  3  2  2  2  2  2  2  2  2  3  2  2  3
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
##  1  2  3  2  2  2  2  3  3  3  3  3  2  3  3  3  3  3  3  3
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
##  2  2  3  3  3  3  3  3  1  3  3  2  2  3  3  3  3  3  3  2
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
##  3  3  3  3  3  3  3  3  1  2  3  1  2  1  1  2  1  3  1  1
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
##  1  3  3  3  1  1  3  1  1  1  1  1  1  1  1  1  1  1  1  1
## 101
##  1
##
## Within cluster sum of squares by cluster:
## [1] 1.440856 0.182348 1.026334
## (between_SS / total_SS = 65.5 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

```
planet_kmeans3
```

```
## K-means clustering with 3 clusters of sizes 14, 34, 53
##
## Cluster means:
##      mass      period      eccen
## 1 0.60560786 0.31606632 0.3953614
## 2 0.16777347 0.11500361 0.5343613
## 3 0.09576256 0.07984122 0.1315524
##
## Clustering vector:
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  3  3  2  3  3  3  3  2  3  3  3  3  3  3  3  3  3  3  3  2
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
##  2  3  2  3  3  3  3  3  2  2  2  2  3  2  3  3  3  3  2  2
## 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60
##  3  3  2  3  3  2  2  2  2  3  3  3  3  2  3  3  3  3  3  3
## 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80
##  3  3  3  2  2  3  3  2  2  3  2  2  3  1  2  3  2  2  2  2
## 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100
##  2  3  3  3  2  1  2  1  2  1  1  2  1  1  1  1  1  1  1  1
## 101
##  1
##
## Within cluster sum of squares by cluster:
## [1] 1.719130 1.787017 1.755694
## (between_SS / total_SS =  57.2 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

```
table(score_kmeans3$cluster)
```

```
##
##  1  2  3
## 26 30 45
```

```
table(planet_kmeans3$cluster)
```

```
##
##  1  2  3
## 14 34 53
```

The results are significantly different from the results in (b).

4.2


```
library(ISLR)
data(Default)
```

(a)

```
## split the sample set
set.seed(1234)
train <- sample(nrow(Default), 0.7 * nrow(Default))
Default.train <- Default[train, ]
Default.validate <- Default[-train, ]

table(Default.train$default)
```

```
##
##   No   Yes
## 6769  231
```

```
table(Default.validate$default)
```

```
##
##   No   Yes
## 2898  102
```

```
## logistic regression
fit.logit <-
  glm(default ~ student + balance + income,
       data = Default.train,
       family = binomial())
summary(fit.logit)
```

```
##
## Call:
## glm(formula = default ~ student + balance + income, family = binomial(),
##      data = Default.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1410  -0.1435  -0.0573  -0.0208   3.6583
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.064e+01  5.867e-01 -18.129  < 2e-16 ***
## studentYes  -7.078e-01  2.747e-01  -2.577  0.00998 **
## balance      5.666e-03  2.752e-04  20.586  < 2e-16 ***
## income      -4.095e-07  9.829e-06  -0.042  0.96677
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
## Null deviance: 2030.3 on 6999 degrees of freedom
## Residual deviance: 1111.6 on 6996 degrees of freedom
## AIC: 1119.6
##
## Number of Fisher Scoring iterations: 8
```

```
prob <- predict(fit.logit, Default.validate, type = "response")

logit.pred <- factor(prob > .5,
                     levels = c(FALSE, TRUE),
                     labels = c("No", "Yes"))
logit.perf <- table(Default.validate$default,
                    logit.pred,
                    dnn = c("Actual", "Predicted"))

## confusion matrix
logit.perf
```

```
##      Predicted
## Actual   No  Yes
##    No 2887   11
##    Yes   65   37
```

```
## validation set error
error <-
  (logit.perf[1, 2] + logit.perf[2, 1]) / (nrow(Default.validate))
error
```

```
## [1] 0.02533333
```

(b)


```
library(rpart)
set.seed(1234)

dtree <- rpart(
  default ~ student + balance + income,
  data = Default.train,
  method = "class",
  parms = list(split = "information")
)

dtree$cpable
```

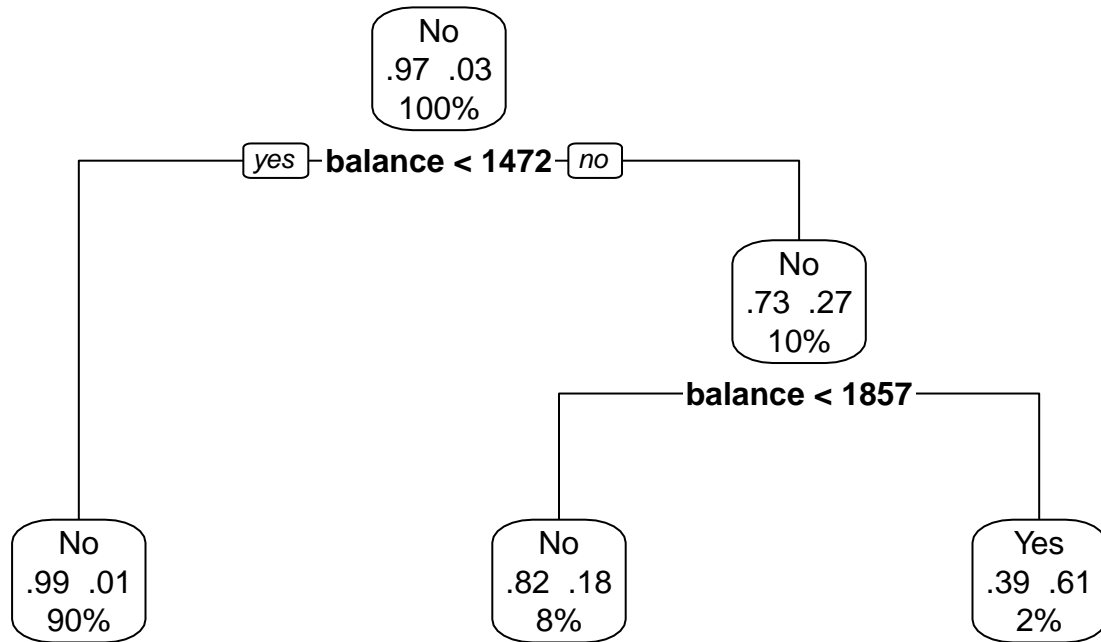
```
##      CP nsplit rel error      xerror      xstd
## 1 0.06926407      0 1.0000000 1.0000000 0.06470044
## 2 0.02597403      2 0.8614719 0.9047619 0.06164232
## 3 0.01298701      4 0.8095238 0.8701299 0.06048665
## 4 0.01010101      6 0.7835498 0.8658009 0.06034044
## 5 0.01000000      9 0.7532468 0.8658009 0.06034044
```

```
plotcp(dtree)
```



```
dtree.pruned <- prune(dtree, cp = 0.039)
library(rpart.plot)
prp(
  dtree.pruned,
  type = 2,
  extra = 104,
  fallen.leaves = TRUE,
  main = "Decision Tree"
)
```

Decision Tree



(C)

```
library(party)
```

```
##      grid
```

```
##      mvtnorm
```

```
##
```

```
##      'mvtnorm'
```

```
## The following object is masked from 'package:mclust':
```

```
##
```

```
##      dmnorm
```

```
##      modeltools
```

```
##      stats4
```

```
##      strucchange
```

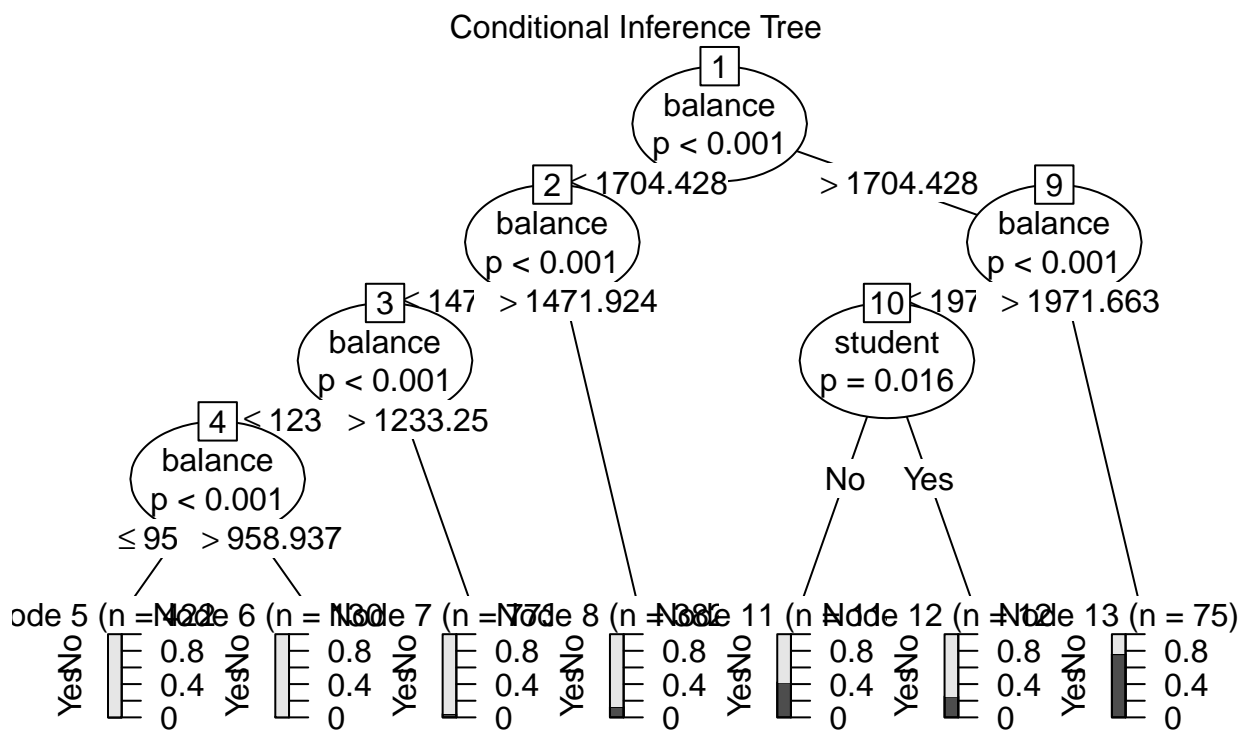
```
##      zoo
```

```
##
## 'zoo'

## The following objects are masked from 'package:base':
##
## as.Date, as.Date.numeric

## sandwich

fit.ctree <- ctree(default ~ student + balance + income,
  data = Default.train)
plot(fit.ctree, main = "Conditional Inference Tree")
```



```
## prediction in the validation set
ctree.pred <-
  predict(fit.ctree, Default.validate, type = "response")

ctree.perf <- table(Default.validate$default,
  ctree.pred,
  dnn = c("Actual", "Predicted"))

## confusion matrix
ctree.perf
```

```
## Predicted
```

```
## Actual    No   Yes
##      No 2888   10
##      Yes  69   33
```

(d)

- traditional decision trees:

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1234)
```

```
## grow the forest
```

```
fit.forest1 <-
  randomForest(
    default ~ student + balance + income,
    data = Default.train,
    na.action = na.roughfix,
    importance = TRUE
  )
fit.forest1
```

```
##
```

```
## Call:
```

```
## randomForest(formula = default ~ student + balance + income, data = Default.train, importance =
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 500
```

```
## No. of variables tried at each split: 1
```

```
##
```

```
##           OOB estimate of error rate: 2.96%
```

```
## Confusion matrix:
```

```
##           No Yes class.error
```

```
## No  6752  17 0.002511449
```

```
## Yes  190  41 0.822510823
```

```
forest1.pred <- predict(fit.forest1, Default.validate)
```

```
forest1.perf <- table(Default.validate$default,
                      forest1.pred,
                      dnn = c("Actual", "Predicted"))
```

```
## confusion matrix
```

```
forest1.perf
```

```
##           Predicted
```

```
## Actual    No   Yes
```

```
##      No 2889    9
```

```
##      Yes  76   26
```

- conditional inference trees:

```
## grow the forest
fit.forest2 <-
  cforest(default ~ student + balance + income,
           data = Default.train,
           controls = cforest_classical(mtry = 2))
fit.forest2

##
## Random Forest using Conditional Inference Trees
##
## Number of trees: 500
##
## Response: default
## Inputs: student, balance, income
## Number of observations: 7000

forest2.pred <- predict(fit.forest2, newdata = Default.validate)
forest2.perf <- table(Default.validate$default,
                     forest2.pred,
                     dnn = c("Actual", "Predicted"))
## confusion matrix
forest2.perf

##      Predicted
## Actual  No  Yes
## No  2887  11
## Yes   65  37
```

Compare the predictive accuracy

```
forest.traditional.accuracy <-
  (forest1.perf[1, 1] + forest1.perf[2, 2]) / nrow(Default.validate)
forest.conditional.accuracy <-
  (forest2.perf[1, 1] + forest2.perf[2, 2]) / nrow(Default.validate)

data.frame(forest.traditional.accuracy, forest.conditional.accuracy)

## forest.traditional.accuracy forest.conditional.accuracy
## 1 0.9716667 0.9746667
```

The accuracy of random forest based on the conditional inference trees is slightly larger than the accuracy of random forest based on the traditional decision trees.

(e)

```
library(e1071)
set.seed(1234)
fit.svm <-
```

```

svm(
  default ~ student + balance + income,
  data = Default.train,
  gamma = 1,
  cost = 1
)
fit.svm

```

```

##
## Call:
## svm(formula = default ~ student + balance + income, data = Default.train,
##      gamma = 1, cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  1
##
## Number of Support Vectors:  616

```

```

svm.pred <- predict(fit.svm, na.omit(Default.validate))
svm.perf <- table(na.omit(Default.validate)$default,
                  svm.pred,
                  dnn = c("Actual", "Predicted"))
svm.perf

```

```

##      Predicted
## Actual   No  Yes
##    No 2891   7
##    Yes  76  26

```

- Compare:

```

performance <- function(table, n = 2) {
  if (!all(dim(table) == c(2, 2)))
    stop("Must be a 2 x 2 table")
  tn <- table[1, 1]
  fp <- table[1, 2]
  fn <- table[2, 1]
  tp <- table[2, 2]
  sensitivity <- tp / (tp + fn)
  specificity <- tn / (tn + fp)
  ppp <- tp / (tp + fp)
  npp <- tn / (tn + fn)
  hitrate <- (tp + tn) / (tp + tn + fp + fn)
  result <-
    data.frame(c(sensitivity,
                  specificity,
                  ppp,
                  npp,
                  hitrate))
}

```



```

    return(result)
}

result <- data.frame(
  svm = performance(svm.perf),
  conditionaltree = performance(ctree.perf),
  forest1 = performance(forest1.perf),
  forest2 = performance(forest2.perf),
  logit = performance(logit.perf),
  row.names = c(
    "sensitivity",
    "specificity",
    "positive predictive power",
    "negative predictive power",
    "accuracy"
  )
)
colnames(result) <-
  c("svm", "cond.tree", "forest.trad", "forest.cond", "logit")
result

```

```

##               svm cond.tree forest.trad forest.cond   logit
## sensitivity      0.2549020 0.3235294   0.2549020   0.3627451 0.3627451
## specificity      0.9975845 0.9965493   0.9968944   0.9962043 0.9962043
## positive predictive power 0.7878788 0.7674419   0.7428571   0.7708333 0.7708333
## negative predictive power 0.9743849 0.9766655   0.9743676   0.9779810 0.9779810
## accuracy        0.9723333 0.9736667   0.9716667   0.9746667 0.9746667

```

Conclusion:

Random forest based on the conditional trees has the same performance as the logistic regression. They have the best sensitivity, negative predictive power, and accuracy.

While svm has the best specificity and positive predictive power.